

Proximity and Data Pre-processing

Huiping Cao

Outline

- Types of data
- Data quality
- Measurement of proximity
- Data pre-process

Similarity and Dissimilarity

- Dissimilarity (distance)
 - Numerical measure of how **different** two data objects are
 - Lower when objects are more alike
 - **Minimum dissimilarity** is often 0
 - **Upper limit varies**
- Similarity
 - Numerical measure of how **alike** two data objects are.
 - Is higher when objects are more alike.
 - Often falls in the **range [0,1]**
- **Proximity** refers to either a similarity or dissimilarity

Similarity and Dissimilarity (cont.)

- Simple attributes, multiple attributes
 - E.g., Dense data: Euclidean distance
 - E.g., Sparse data: Jaccard and Cosine similarity

Similarity/Dissimilarity for Simple Attributes

- p and q are the attribute values for two data objects.

| Attribute type | Dissimilarity | Similarity |
|-------------------|--|---|
| Nominal | $d = \begin{cases} 0 & \text{if } p = q \\ 1 & \text{if } p \neq q \end{cases}$ | $s = \begin{cases} 1 & \text{if } p = q \\ 0 & \text{if } p \neq q \end{cases}$ |
| Ordinal | $d = \frac{ p-q }{n-1}$ (values mapped to integers 0 to $n-1$, where n is the number of values) | $s = 1 - \frac{ p-q }{n-1}$ |
| Interval or Ratio | $d = p - q $ | $s = -d, s = \frac{1}{d+1}$ or $s = 1 - \frac{d - \min_d}{\max_d - \min_d}$ |

where \min_d and \max_d are the minimum and maximum distances between every two values.

Similarity and Dissimilarity – multiple attributes


- Euclidean, Minkowski, DTW, etc.
- SMC, Jaccard, Cosine, Hamming, etc.

Euclidean Distance

$$dist = \sqrt{\sum_{i=1}^n (q_i - q'_i)^2}$$

- n is the number of dimensions (attributes)
- q_i and q'_i : the i th attributes (components) of data objects q and q' .
- Standardization is necessary, if scales differ.

Euclidean Distance – Example



| point | x | y |
|-------|---|---|
| p1 | 0 | 2 |
| p2 | 2 | 0 |
| p3 | 3 | 1 |
| p4 | 5 | 1 |

| | p1 | p2 | p3 | p4 |
|----|-------|-------|-------|-------|
| p1 | 0 | 2.828 | 3.162 | 5.099 |
| p2 | 2.828 | 0 | 1.414 | 3.162 |
| p3 | 3.162 | 1.414 | 0 | 2 |
| p4 | 5.099 | 3.162 | 2 | 0 |

Distance Matrix

Manhattan Distance

- Considered by Hermann Minkowski in 19th century Germany.

$$dist = \left(\sum_{i=1}^n |q_i - q'_i| \right)$$

- Taxi cab metric/distance, rectilinear distance, Minkowski's L_1 norm/distance, city block distance, or Manhattan length.
- Parameters
 - n is the number of dimensions (attributes)
 - q_i and q'_i are, respectively, the i th attributes (components) of data objects q and q' .
- Applications
 - In chess, the distance between squares on the chessboard for rooks is measured in Manhattan distance.
 - The length of the shortest path a taxi could take between two intersections equal to the intersections' distance in taxicab geometry.
 - Integrated circuits where wires only run parallel to the X or Y axis.

Minkowski Distance

- The Minkowski distance is a metric on Euclidean space which can be considered as a generalization of both the Euclidean distance and the Manhattan distance.

$$dist = \left(\sum_{i=1}^n |q_i - q'_i|^p \right)^{\frac{1}{p}}$$

- Parameters
 - p is a parameter
 - n is the number of dimensions (attributes)
 - q_i and q'_i are, respectively, the i th attributes (components) of data objects q and q' .

Minkowski Distance: Illustration

- $p = 1$. Manhattan distance.
- $p = 2$. Euclidean distance.
- $p \rightarrow \infty$. Supremum, Chebyshev (L_{max} norm, L_∞ norm) distance.

$$\lim_{p \rightarrow \infty} \left(\sum_{i=1}^n |q_i - q'_i|^p \right)^{\frac{1}{p}} = \max_{i=1}^n |q_i - q'_i|$$

Kings and queens use Chebyshev distance in Chess.

- L_{min} norm

$$\lim_{p \rightarrow -\infty} \left(\sum_{i=1}^n |q_i - q'_i|^p \right)^{\frac{1}{p}} = \min_{i=1}^n |q_i - q'_i|$$

Note: Do not confuse p with n , i.e., all these distances are defined for all numbers of dimensions.

Minkowski Distance– Example

| point | x | y |
|-------|---|---|
| p1 | 0 | 2 |
| p2 | 2 | 0 |
| p3 | 3 | 1 |
| p4 | 5 | 1 |

| L1 | p1 | p2 | p3 | p4 |
|----|----|----|----|----|
| p1 | 0 | 4 | 4 | 6 |
| p2 | 4 | 0 | 2 | 4 |
| p3 | 4 | 2 | 0 | 2 |
| p4 | 6 | 4 | 2 | 0 |

| L2 | p1 | p2 | p3 | p4 |
|----|-------|-------|-------|-------|
| p1 | 0 | 2.828 | 3.162 | 5.099 |
| p2 | 2.828 | 0 | 1.414 | 3.162 |
| p3 | 3.162 | 1.414 | 0 | 2 |
| p4 | 5.099 | 3.162 | 2 | 0 |

| L_∞ | p1 | p2 | p3 | p4 |
|------------|----|----|----|----|
| p1 | 0 | 2 | 3 | 5 |
| p2 | 2 | 0 | 1 | 3 |
| p3 | 3 | 1 | 0 | 2 |
| p4 | 5 | 3 | 2 | 0 |

Distance Matrix

Issues to Consider in Calculating Distance

- Attributes have **different scales**
- Attributes are **correlated**
- Objects are composed of **different types** of attributes (Qualitative vs. quantitative)
- Attributes have **different weights**

Common Properties of a Distance

- Distances, such as the Euclidean distance, have some well known properties.
 - **Positivity:**
 $d(p, q) \geq 0$ for all p and q ;
 $d(p, q) = 0$ only if $p = q$.
 - **Symmetry:** $d(p, q) = d(q, p)$ for all p and q .
 - **Triangle Inequality:** $d(p, r) \leq d(p, q) + d(q, r)$ for all points p , q , and r .
- A distance that satisfies these properties is a **metric**.

Elastic distances

- Dynamic Time Warping (DTW)
- Edit distance based measure
 - Longest Common SubSequence (LCSS)
 - Edit Distance on Real Sequence (EDR)

Elastic distances - DTW

- In time series analysis, **Dynamic Time Warping (DTW)** is an algorithm for measuring similarity between two temporal sequences which may vary in time or speed.
- **DTW** is a technique to find an **optimal alignment**¹ between time series if one time series can be “warped” non-linearly by stretching or shrinking it.
- **Applications**: similarities in walking patterns, video, audio, etc.
- Calculates an **optimal match** between two given sequences (e.g., time series) with certain restrictions

¹M. Müller, Information Retrieval for Music and Motion, Springer, 2007, ISBN:

978-3-540-74047-6, <http://www.springer.com/978-3-540-74047-6>

Elastic distances - DTW - Problem definition

- Given a **reference sequence** $s[1, \dots, n]$ and a **query sequence** $t[1, \dots, m]$ and local distance measure $dist(s[i], t[j])$
- **Goal**: find an alignment between s and t having minimal overall cost
- An **(n, m) -warping path** (or simply referred to as **warping path** if n and m are clear from the context) is a sequence $p = (p_1, \dots, p_L)$ with $p_i = (n_i, m_i) \in [1 : n] \times [1 : m]$ for i ($1 \leq i \leq L$) satisfying the following three conditions
 - **Boundary condition**: $p_1 = (1, 1)$ and $p_L = (n, m)$
 - **Monotonicity condition**: $n_1 \leq n_2 \leq \dots \leq n_L$ and $m_1 \leq m_2 \leq \dots \leq m_L$
 - **Step size condition**: $p_{i+1} - p_i$ is in $\{(1, 0), (0, 1), (1, 1)\}$ for i ($1 \leq i \leq L - 1$)

Elastic distances - DTW - Distance

- Total cost of a warping path $p=(p_1, \dots, p_L)$ between s and t is

$$cost_p(s, t) = \sum_{i=1}^L dist(p_i) = \sum_{i=1}^L dist(s[n_i], t[m_i])$$

where $p_i = (n_i, m_i)$.

- Let p^* be an optimal warping path between s and t
- The DTW distance $DTW(s, t)$ between s and t is defined as the total cost of p^* .

$$\begin{aligned} DTW(s, t) &= cost_{p^*}(s, t) \\ &= \min \{ cost_p(s, t) | p \text{ is an } (n, m)\text{-warping path} \} \end{aligned}$$

Check [algorithm](#).

DTW calculation using Python

Reference

```
import numpy as np
from fastdtw import fastdtw
from scipy.spatial.distance import euclidean

x = np.array([1, 2, 3, 3, 7])
y = np.array([1, 2, 2, 2, 2, 2, 2, 4])

distance, path = fastdtw(x, y, dist=euclidean)

print(distance)
print(path)

# 5.0
# [(0, 0), (1, 1), (1, 2), (1, 3), (1, 4), (2, 5), (3, 6), (4, 7)]
```

Another good library:

<https://pyts.readthedocs.io/en/stable/generated/pyts.metrics.dtw.html>

Common Properties of a Similarity

- Similarities also have some well known properties.
 - **Maximum similarity:** $s(p, q) = 1$ only if $p = q$
 - **Symmetry:** $s(p, q) = s(q, p)$ for all p and q
- $s(p, q)$ is the similarity between points (data objects) p and q .

Similarity Between Binary Vectors

- Compute similarities using the following quantities
 - M_{01} = the number of attributes where p was 0 and q was 1
 - M_{10} = the number of attributes where p was 1 and q was 0
 - M_{00} = the number of attributes where p was 0 and q was 0
 - M_{11} = the number of attributes where p was 1 and q was 1
- Similarity coefficient
- In the range of $[0,1]$

Distance

- Simple Matching Coefficient (SMC)
- Jaccard Coefficient
- Cosine Similarity
- Hamming distance
- Correlation
- Mutual information (MI)

Simple Matching Coefficient (SMC)

$$SMC = \frac{\text{number of matches}}{\text{number of attributes}} = \frac{M_{11} + M_{00}}{M_{01} + M_{10} + M_{11} + M_{00}}$$

Example:

- $p = 1\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0$
- $q = 0\ 0\ 0\ 0\ 0\ 0\ 1\ 0\ 0\ 1$
- $M_{01} = 2, M_{10} = 1, M_{00} = 7, M_{11} = 0$
- $SMC = \frac{7+0}{10} = 0.7$

Jaccard Coefficient

$$J = \frac{\# \text{ of 11 matches}}{\text{total \# of attributes} - \# \text{ of 00 matches}} = \frac{M_{11}}{M_{01} + M_{10} + M_{11}}$$

- Handle asymmetric binary attributes
- Example:

- $p = 1\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0$

- $q = 0\ 0\ 0\ 0\ 0\ 0\ 1\ 0\ 0\ 1$

- $M_{01} = 2, M_{10} = 1, M_{00} = 7, M_{11} = 0$

- $J = \frac{0}{2+1+0} = 0$

Cosine Similarity

- If d_1 and d_2 are two document vectors, then

$$\cos(d_1, d_2) = \frac{d_1 \cdot d_2}{||d_1|| ||d_2||}$$

- \cdot indicates vector dot product
- $||d||$ is the length of vector d
- $d_1 = 3 \ 2 \ 0 \ 5 \ 0 \ 0 \ 0 \ 2 \ 0 \ 0$
- $d_2 = 1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 1 \ 1 \ 0 \ 2$
- $d_1 \cdot d_2 = 3 + 2 = 5$
- $||d_1|| = \sqrt{42} = 6.481$
- $||d_2|| = \sqrt{7} = 2.646$
- $\cos(d_1, d_2) = \frac{5}{6.481 \cdot 2.646} = 0.2916$

Hamming distance

- Measure distance among vectors of **equal length** with **nominal** values.
- The number of positions at which the corresponding symbols are **different**.
- Example
 - “toned” and “roses” is 3.
 - 1011101 and 1001001 is 2.
 - 2173896 and 2233796 is 3.
- Named after *Richard Hamming*. It is used in telecommunication to **count the number of flipped bits** in a fixed-length binary word as an estimate of error, and therefore is sometimes called the **signal distance**.
- Not suitable for comparing strings of **different lengths**, or strings where not just substitutions but also insertions or deletions have to be expected

Correlation

- Correlation measures the **linear relationship** between objects
- To compute correlation, we standardize data objects, p and q , and then take their dot product

$$\text{corr}(p, q) = \frac{\text{covariance}(p, q)}{\text{std}(p) * \text{std}(q)} = \frac{s_{pq}}{s_p * s_q}$$

$$\text{covariance}(p, q) = s_{pq} = \frac{1}{n-1} \sum_{i=1}^n (p_i - \bar{p})(q_i - \bar{q})$$

$$\text{std}(p) = s_p = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (p_i - \bar{p})^2}$$

$$\text{std}(q) = s_q = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (q_i - \bar{q})^2}$$

$$\bar{p} = \frac{1}{n} \sum_{i=1}^n p_i, \bar{q} = \frac{1}{n} \sum_{i=1}^n q_i$$

Correlation - Example 1

$$x = (-3, 6, 0, 3, -6)$$

$$y = (1, -2, 0, -1, 2)$$

Compute the Pearson's correlation

$$\text{corr}(x, y) = -1, x = -3y$$

Correlation - Example 2

$$x = (-3, -2, -1, 0, 1, 2, 3)$$

$$y = (9, 4, 1, 0, 1, 4, 9)$$

- Compute the Pearson's correlation $\text{corr}(x, y) = 0$, $y = x^2$
- No **linear** relationship, but **nonlinear** relationships can still exist.
- Pearson's correlation only measures **linear** correlation.

Correlation - Python Example

```
from scipy.stats import pearsonr

p = np.array([1,3,5])
q = np.array([2,3,5])

corr, _ = pearsonr(p, q)
print('Pearsons correlation: %.3f' % corr)

#Pearsons correlation: 0.982
```

Mutual information

- For measuring nonlinear correlation
- Come from information theory

$$I(X, Y) = H(X) + H(Y) - H(X, Y)$$

where

- $H(X) = -\sum_{j=1}^m P(X = u_j) \log_2 P(X = u_j)$, the entropy (average information) of X
- $H(Y) = -\sum_{k=1}^n P(Y = v_k) \log_2 P(Y = v_k)$, the entropy (average information) of Y
- $H(X, Y) = -\sum_{j=1}^m \sum_{k=1}^n P(X = u_j, Y = v_k) \log_2 P(X = u_j, Y = v_k)$, the joint entropy of X and Y
- Minimum value is 0, representing that the value of one variable tells us nothing about the another.
- Maximum value (no fixed value), representing that one variable completely depend on another.

Mutual information - Example

$x = (-3, -2, -1, 0, 1, 2, 3)$

$y = (9, 4, 1, 0, 1, 4, 9)$

- Compute the mutual information $I(X, Y) = 1.9502$

General Approach for Combining Similarities

- Situation: attributes are of **different types**, an overall similarity is needed
- For the i th attribute, compute a similarity $s_i \in [0, 1]$
- Define δ_i for the i th attribute
 - $\delta_i = 0$ (i.e., do not count this attribute)
 - if the i th attribute is an asymmetric attribute and both objects have a value of 0
 - or if one of the objects has a missing value for the i th attribute
 - $\delta_i = 1$ otherwise
- Compute the overall similarity

$$\text{similarity}(p, q) = \frac{\sum_{i=1}^n \delta_i s_i}{\sum_{i=1}^n \delta_i}$$

Using Weights to Combine Similarities

- May not want to treat all attributes the same.
- Use weights $w_i \in [0, 1]$ and $\sum_{i=1}^n w_i = 1$.

$$\text{similarity}(p, q) = \frac{\sum_{i=1}^n w_i \delta_i s_i}{\sum_{i=1}^n \delta_i}$$

Sampling

- It is often used for both the preliminary investigation of the data and the final data analysis.
- Statisticians sample because obtaining the entire set of data of interest is too expensive or time consuming.
- Sampling is used in data mining because processing the entire set of data of interest is too expensive or time consuming.


Key Principle for Effective Sampling

- Using a sample will work almost as well as using the entire data sets, if the sample is representative
- A sample is **representative** if it has approximately the same property (of interest) as the original set of data


Types of Sampling

- **Simple Random Sampling:** there is an equal probability of selecting any particular item
- **Variations**
 - **Sampling without replacement:** As each item is selected, it is removed from the population
 - **Sampling with replacement:** Objects are not removed from the population as they are selected for the sample.
 - In sampling with replacement, the same object can be picked up more than once
- **Stratified sampling:** Split the data into several partitions; then draw random samples from each partition

Sample Size



8000 points



2000 Points



500 Points

Curse of Dimensionality

- Data analysis becomes **significantly harder** as the dimensionality of the data **increases**.
- Data becomes **increasingly sparse** in the space that it occupies

Dimensionality Reduction

■ Purpose:

- Avoid curse of dimensionality
- Reduce amount of time and memory required by data mining algorithms
- Allow data to be more easily visualized
- May help to eliminate irrelevant features or reduce noise

■ Techniques

- Principle Component Analysis (PCA)
- Singular Value Decomposition (SVD)
- Others: supervised and non-linear techniques

Feature Subset Selection

- Another way to **reduce dimensionality** of data
- **Redundant** features
 - Duplicate much or all of the information contained in one or more other attributes
 - Example: purchase price of a product and the amount of sales tax paid
- **Irrelevant** features
 - Contain no information that is useful for the data mining task at hand
 - Example: students' ID is often irrelevant to the task of predicting students' GPA

Feature Subset Selection–Techniques (S.S.)

- **Brute-force** approach:
 - Try all possible feature subsets as input to data mining algorithm
- **Embedded** approaches:
 - Feature selection occurs naturally as part of the data mining algorithm
- **Filter** approaches
 - Features are selected before data mining algorithm is run
- **Wrapper** approaches:
 - Use the data mining algorithm as a black box to find best subset of attributes

Feature Creation (S.S.)

- Create **new attributes** that can capture the important information in a data set much **more efficiently** than the original attributes
- Three general methodologies
 - **Feature extraction**: domain-specific
 - **Mapping data to new space**
 - Fourier transform
 - Wavelet transform
 - **Feature construction**: **combining features**. E.g., Density in cluster analysis.

- Chapter 2: Introduction to Data Mining (2nd Edition) by Pang-Ning Tan, Michael Steinbach, Anuj Karpatne, and Vipin Kumar.
- Textbook website:
`https://www-users.cse.umn.edu/~kumar001/dmbook/index.php`
- `https://scikit-learn.org/stable/modules/classes.html#module-sklearn.metrics.pairwise`